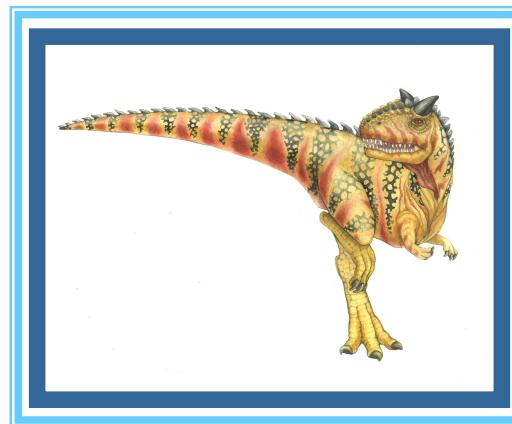


Chapter 5: Process Synchronization





Discussed till now

- Race condition
 - Demonstration through Producer-Consumer problem
- Critical Section Problem
 - Requirements for solutions to the problem
- Solutions to Critical Section Problem
 - Software solution – Peterson's algorithm
 - Hardware solutions (special atomic instructions) -
 - ▶ `test_and_set()`
 - ▶ `compare_and_swap()`





Mutex Locks





Mutex Locks

- Previous hardware-based solutions are complicated and generally inaccessible to application programmers
- OS designers build **software tools** to solve critical section problem
- Simplest such tool is mutex lock
- Protect a critical section by first **acquire()** a lock then **release()** the lock
 - Boolean variable indicating if lock is available or not
- **Calls to `acquire()` and `release()` must be atomic**
 - Usually implemented via hardware atomic instructions





acquire() and release()

- ```
acquire() {
 while (!available)
 ; /* busy wait */
 available = false;
}

■ release() {
 available = true;
}

■ do {
 acquire lock
 critical section
 release lock
 remainder section
} while (true);
```





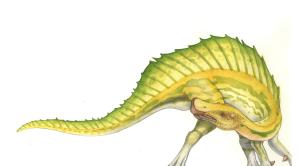
# Mutex Locks – adv and disadv

## ■ Disadvantage

- Requires **busy waiting** (wastage of CPU cycles that some other process could have used)
- This type of mutex lock is also called a **spinlock** (process "spins" while waiting for the lock to become available)

## ■ Advantages

- No context switch when a process must wait on a lock
- Useful when locks are expected to be held for short time intervals
- Can be employed on multiprocessor systems where one thread can "spin" on one processor while another thread performs its critical section on another processor





# Mutex Locks for threads

---

```
int pthread_mutex_init(pthread_mutex_t *mutex,
 const pthread_mutexattr_t *attr);

int pthread_mutex_lock(pthread_mutex_t *mutex);

int pthread_mutex_unlock(pthread_mutex_t *mutex);

int pthread_mutex_destroy(pthread_mutex_t *mutex);
```





# Mutex Locks for threads (contd.)

```
pthread_mutex_t lock; // global variable

if (pthread_mutex_init(&lock, NULL) != 0) {
 printf("Error");
 return 1;
}

...
...
pthread_mutex_lock(&lock);
<critical section>
pthread_mutex_unlock(&lock);

...
...
pthread_mutex_destroy(&lock);
```





# Semaphores





# Semaphore

- Synchronization tool that provides more sophisticated ways (than Mutex locks) for process to synchronize their activities.
- Semaphore **s** – integer variable
- Can only be accessed via two indivisible / atomic operations (apart from initialization)
  - **wait()** and **signal()**
  - Originally called **P()** and **V()**
  - P() possibly from "proberen" (to test)
  - V() possibly from "verhogen" (to increase)





# Semaphore operations

- Definition of the wait( ) operation

```
wait(s) {
 while (s <= 0)
 ; // busy wait
 s--;
}
```

- Definition of the signal( ) operation

```
signal(s) {
 s++;
}
```

- Note: both these operations must be **atomic** (executed indivisibly)





# Types of Semaphore

- **Counting semaphore** – integer value can range over an unrestricted domain
- **Binary semaphore** – integer value can range only between 0 and 1
  - Same as a **mutex lock**





# Semaphore Usage

- Can solve various synchronization problems
  - Mutual exclusion
  - Process synchronization
  - Resource counting
- Use a binary semaphore **s** (value can be 0 or 1)
  - Initialize **s** to 1
  - Every cooperating process has this structure:

**P<sub>i</sub>:**

```
wait(s) ;
<Critical section>
signal(s) ;
<Remainder section>
```





# Semaphore Usage

- Can solve various synchronization problems
  - Mutual exclusion
  - Process synchronization
  - Resource counting
- Consider  $P_1$  (containing statement  $S_1$ ) and  $P_2$  (containing statement  $S_2$ ) that require  $S_1$  to happen before  $S_2$   
Create a semaphore “synch” initialized to 0

**P1:**

```
S1;
 signal(synch);
```

**P2:**

```
wait(synch);
S2;
```





# Semaphore Usage

- Can solve various synchronization problems
  - Mutual exclusion
  - Process synchronization
  - **Resource counting**
- We want to keep track of a given resource with a finite number of instances
- Use a counting semaphore whose value can range over an unrestricted domain
  - Initialize to the number of instances available
  - Each process that wishes to use an instance of the resource performs `wait()`
  - When a process releases its instance of the resource, perform `signal()`
- Can this be used to solve a problem we have discussed earlier?





# Semaphore Implementation

---

- Must guarantee that no two processes can execute the **wait()** and **signal()** on the same semaphore at the same time
- Applications may spend lots of time in critical sections; hence important to avoid busy waiting in the implementation of **wait()** and **signal()**





# Avoiding busy waiting

- With each semaphore there is an associated waiting queue
- If a process executes **wait(s)** and finds that the semaphore **s** value is not positive, process blocks itself
  - Place the process into the waiting queue for **s**
  - Control transfer to CPU scheduler for selecting next process to execute
- If a process executes **signal(s)**
  - A process that is blocked on **s** should be restarted
  - Move this process from waiting queue to ready queue





# Semaphore Implementation with no busy waiting

- With each semaphore there is an associated waiting queue
- Each entry in a waiting queue has two data items:
  - value (of type integer)
  - pointer to next record in the list
- Two operations:
  - **block** – place the process invoking the operation on the appropriate waiting queue
  - **wakeup** – remove one of processes in the waiting queue and place it in the ready queue
- ```
typedef struct{
    int value;
    struct process *list;
} semaphore;
```





Implementation with no Busy waiting (Cont.)

```
wait(semaphore *S) {  
    S->value--;  
    if (S->value < 0) {  
        add this process to S->list;  
        block();  
    }  
}  
  
signal(semaphore *S) {  
    S->value++;  
    if (S->value <= 0) {  
        remove a process P from S->list;  
        wakeup(P);  
    }  
}
```

The block() and wakeup() operations must be provided by the kernel as basic system calls





Some observations

- The list of processes waiting for various semaphores can be easily implemented using a link field in the PCB
 - E.g., as a FIFO queue
- The wait() and signal() operations must execute atomically
 - Disabling interrupts can be a solution for uniprocessor systems (but difficult for multiprocessor systems)
 - Alternative locking via hardware instructions such as TAS or CAS can be used
- Can the value of a counting semaphore become negative?





Deadlock and Starvation

- **Deadlock** – two or more processes are waiting indefinitely for an event that can be caused only by one of the waiting processes
- Let S and Q be two semaphores initialized to 1

P_0	P_1
<code>wait(S);</code>	<code>wait(Q);</code>
<code>wait(Q);</code>	<code>wait(S);</code>
...	...
<code>signal(S);</code>	<code>signal(Q);</code>
<code>signal(Q);</code>	<code>signal(S);</code>

- **Starvation – indefinite blocking**

- A process may never be removed from the semaphore queue in which it is suspended





Problems with Semaphores

- Incorrect use of semaphore operations:
 - signal (mutex) wait (mutex)
 - wait (mutex) ... wait (mutex)
 - Omitting of wait (mutex) or signal (mutex) (or both)
- Deadlock and starvation are possible.





POSIX Semaphores

- For using semaphores
 - `#include semaphore.h`
 - Compile the code by linking with `-lpthread -lrt`
- Declare a semaphore: `sem_t mutex;`
- Initialize a semaphore:
 - `sem_init (sem_t *sem, int pshared, unsigned int value)`
 - **sem** - the semaphore to be initialized
 - **pshared** - whether or not the newly initialized semaphore is shared between processes or between threads. Non-zero => the semaphore can be accessed by other processes. Zero => can be shared only by threads belonging to this process
 - **value** - the value to assign to the newly initialized semaphore.
 - E.g., `sem_init(&mutex, 0, 1); // initialize to 1`





POSIX Semaphores (contd.)

- Wait on a semaphore
 - `int sem_wait (sem_t *sem)`
 - E.g., `sem_wait(&mutex);`
- Signal a semaphore
 - `int sem_post (sem_t *sem)`
 - E.g., `sem_post(&mutex);`
- Destroy a sempahore
 - `sem_destroy(sem_t *sem)`
 - E.g., `sem_destroy(&mutex);`





Classical Problems of Synchronization

- Bounded-Buffer Producer-Consumer Problem
- Dining-Philosophers Problem
- Readers and Writers Problem





Bounded-Buffer Problem

- Producer and consumer processes share the following data:
 - n (number of buffers, each of which can hold one item)
 - Semaphore **mutex** initialized to the value 1
 - Semaphore **full** initialized to the value 0
 - Semaphore **empty** initialized to the value n
- Semaphores **empty** and **full** count the number of empty and full buffers
- Semaphore **mutex** provides mutual exclusion in accessing buffer





Bounded Buffer Problem (Cont.)

■ The structure of the **producer** process

```
do {  
    ...  
    /* produce an item in next_produced */  
    ...  
    wait(empty);  
    wait(mutex);  
    ...  
    /* add next produced to the buffer */  
    ...  
    signal(mutex);  
    signal(full);  
} while (true);
```





Bounded Buffer Problem (Cont.)

■ The structure of the **consumer** process

```
do {  
    wait(full);  
    wait(mutex);  
    ...  
    /* remove an item from buffer to next_consumed */  
    ...  
    signal(mutex);  
    signal(empty);  
    ...  
    /* consume the item in next consumed */  
    ...  
} while (true);
```





Dining-Philosophers Problem



- Philosophers spend their lives alternating thinking and eating
 - Don't interact with their neighbors
 - When hungry, try to pick up 2 chopsticks (one at a time) to eat from bowl
 - Need both chopsticks to eat, then release both when done
- In the case of 5 philosophers and 5 single chopsticks
 - Shared data
 - ▶ Semaphore **chopstick [5]** all initialized to 1





Dining-Philosophers Problem Algorithm

■ The structure of Philosopher *i*:

```
do {  
    wait (chopstick[i]);  
    wait (chopStick[(i + 1) % 5]);  
  
    /* eat for a while */  
  
    signal (chopstick[i]);  
    signal (chopstick[(i + 1) % 5]);  
  
    /* think for a while */  
  
} while (TRUE);
```

■ What is the problem with this algorithm?





Dining-Philosophers Problem Algorithm (Cont.)

- Can cause deadlock --> all philosophers can starve to death
- Deadlock handling
 - Allow at most 4 philosophers to be sitting simultaneously at the table
 - Allow a philosopher to pick up the chopsticks only if both are available
 - Use an asymmetric solution -- an odd-numbered philosopher picks up first the left chopstick and then the right chopstick. Even-numbered philosopher picks up first the right chopstick and then the left chopstick





Readers-Writers Problem

- A data set is shared among a number of concurrent processes
 - Readers – only read the data set; they do *not* perform any updates
 - Writers – can both read and write

- Problem –
 - Multiple readers can be allowed to read at the same time
 - A writer must have exclusive access to the shared data while writing





Readers-Writers Problem Variations

- Several variations of how readers and writers are considered – all involve some form of priorities
- **First variation** – no reader is kept waiting unless writer has already obtained permission to use shared dataset
 - No reader should wait for other readers to finish, simply because a writer is waiting
 - Readers have priority, Writers may starve
- **Second** variation – once writer is ready, it performs the write ASAP
 - Writer has priority
- Both may have starvation leading to even more variations
- **We discuss a solution to the first readers-writers problem**





Readers-Writers Problem

- Readers and Writers share the following data:
 - Data set
 - Semaphore **rw_mutex** initialized to 1
 - Semaphore **mutex** initialized to 1
 - Integer **read_count** initialized to 0
- Semaphore **rw_mutex** used as mutual exclusion semaphore for writers
- Semaphore **mutex** is used to ensure mutual exclusion when **read_count** is updated





Readers-Writers Problem (Cont.)

- The structure of a **writer** process

```
do {  
    wait(rw_mutex);  
    ...  
    /* writing is performed */  
    ...  
    signal(rw_mutex);  
} while (true);
```





Readers-Writers Problem (Cont.)

■ The structure of a **reader** process

```
do {  
    wait(mutex);  
    read_count++;  
    if (read_count == 1)  
        wait(rw_mutex);  
    signal(mutex);  
    ...  
    /* reading is performed */  
    ...  
    wait(mutex);  
    read_count--;  
    if (read_count == 0)  
        signal(rw_mutex);  
    signal(mutex);  
} while (true);
```

If a writer is in the critical section, and n readers are waiting:

Which readers are queued up on which semaphore?





Readers-Writers Problem (Cont.)

■ The structure of a **reader** process

```
do {  
    wait(mutex);  
    read_count++;  
    if (read_count == 1)  
        wait(rw_mutex);  
    signal(mutex);  
    ...  
    /* reading is performed */  
    ...  
    wait(mutex);  
    read_count--;  
    if (read_count == 0)  
        signal(rw_mutex);  
    signal(mutex);  
} while (true);
```

If a writer is in the critical section, and n readers are waiting:

- one reader is queued on **rw_mutex**
- the rest $n-1$ readers are queued on **mutex**





Summary

- Race condition
 - Demonstration through Producer-Consumer problem
- Critical Section Problem
 - Requirements for solutions to the problem
- Solutions to Critical Section Problem
 - Software solution – Peterson's algorithm
 - Hardware solutions (special atomic instructions) -
 - ▶ `test_and_set()`
 - ▶ `compare_and_swap()`
 - Mutex locks
 - Semaphores
 - ▶ Handling classical problems of synchronization
 - ▶ POSIX semaphores

